

# 1.6 Floating-point numbers

## Floating-point numbers and normalized scientific notation

An **integer** is a whole number, like 42, 0, or -95. A **floating-point number** is a real number, like 98.6, 0.0001, or -666.667. The "point" refers to the decimal point being able to appear anywhere ("float") in the number.

- Integers are typically used for values that can be counted, like 42 cars, 0 pizzas, or -95 days.
- Floating-point numbers are typically used for values that are measured, like 98.6 degrees, 0.00001 meters, or -666.66

To improve readability and consistency, floating-point numbers are commonly written using **normalized scientific notation**,  $10^1$ ,  $1.0 \times 10^{-4}$ , or  $-6.66667 \times 10^2$ , where the number is written as a digit (+/- 1 to 9), decimal point, fractional part, times  $10^{\text{exponent}}$ . The term "normalized" is in contrast to non-normalized where more than one digit, or a 0, may precede the decimal point, such as  $0.1$  or  $0.1 \times 10^{-3}$ .

The parts of scientific notation are named **significand** for the part before  $\times$  and **exponent** for the power of 10: significand  $\times 10^{\text{exponent}}$ . If the exponent is 0, the power of ten part is sometimes omitted, as in 5.7.

In binary, normalized scientific notation consists of  $1.f \times 2^{\text{exponent}}$ , like  $1.010 \times 2^5$ .  $f$  is the fractional part.

### PARTICIPATION ACTIVITY

#### 1.6.1: Normalized scientific notation: Decimal.

Indicate which numbers are in decimal normalized scientific notation.

1)  $2.05 \times 10^3$

☐ Yes

☐ No

2)  $0.50 \times 10^3$

☐

☐ Yes

☐ No

3)  $27.8 \times 10^3$

☐ Yes

☐ No

4) 3.5

☐ Yes

☐ No

5)  $-5.77 \times 10^3$

☐ Yes

☐ No

6)  $2.05 \times 10^{-3}$

☐ Yes

☐ No

7) 0.0

☐ Yes

☐ No

**PARTICIPATION  
ACTIVITY**

1.6.2: Normalized scientific notation: Binary.

Indicate which numbers are in binary normalized scientific notation.

1) 0.0

☐ Yes

☐ No

2)  $1.01 \times 2^3$

☐ Yes

☐ No

3)  $0.10 \times 2^3$

☐ Yes

☐ No

4)  $11.10 \times 2^3$

☐ Yes

☐ No

5)  $-1.01 \times 2^3$

☐ Yes

☐ No

6)  $1.01 \times 2^{-3}$

☐ Yes

☐ No

**CHALLENGE  
ACTIVITY**

1.6.1: Normalized scientific notation.

Start

Write the number in normalized scientific notation.

Use ^ for exponents. Ex:  $10^4$  for **10<sup>4</sup>**.

177.757 =

3

4

1 2 3 4

Check Next

## Fractions in binary

Fractions in binary are similar to fractions in decimal. In decimal, each digit after the decimal point has weight  $1/10^1$ ,  $1/10^2$ , etc. In binary, each digit after the binary point has weight  $1/2^1$ ,  $1/2^2$ ,  $1/2^3$ ,  $1/2^4$ , etc. (so  $1/2$ ,  $1/4$ ,  $1/8$ ,  $1/16$ , etc.). Ex:  $1.1101$  is  $1 + 1/2 + 1/4$ . Note that for binary numbers, the "dot" is called a **binary point** (versus decimal point for decimal numbers). The general term

Figure 1.6.1: Fractional digit weights for decimal and binary.

Decimal	—	—	—	•	—	—	—	—
	$10^2$	$10^1$	$10^0$		$10^{-1}$	$10^{-2}$	$10^{-3}$	$10^{-4}$
	100	10	1		1/10	1/100	1/1000	1/10000

  

Binary	—	—	—	•	—	—	—	—	Ex: $1.01 = 1 + 1/4 = 1.25$
	$2^2$	$2^1$	$2^0$		$2^{-1}$	$2^{-2}$	$2^{-3}$	$2^{-4}$	
	4	2	1		1/2	1/4	1/8	1/16	
					0.5	0.25	0.125	0.0625	

Converting decimal to binary or vice-versa when fractions are involved uses the same process as without fractions. However, manually converting a decimal with a fraction into binary, converting the whole and fraction parts separately, then concatenating is easier. Ex: For 12.25, 12 is  $1100_2$ , and 0.25 is  $0.01_2$ , yielding  $1100.01$ .

If a decimal fraction cannot be exactly represented as a binary fraction using limited bits, one gets as close as possible, over and over. To represent 0.8 with only four binary fraction bits:

- 0.1100 is  $1/2 + 1/4 = 0.75$ , which is under by 0.05.
- 0.1110 is 0.875, which is over by 0.075.
- 0.1101 is 0.8125, also over but only by 0.0125, which is closest and so is the best representation.

Obviously, more bits means binary fraction values can be closer.

**PARTICIPATION  
ACTIVITY**

## 1.6.3: Fractions in binary.

1) 1.1 binary = \_\_\_\_ decimal. Type as: #.#

**Check**[Show answer](#)

2) 0.001 binary = \_\_\_\_ decimal. Type as:  
0.###

**Check**[Show answer](#)

3) 10.101 binary = \_\_\_\_ decimal. Type as:  
#.###

**Check**[Show answer](#)

4) 0.75 decimal = \_\_\_\_ binary. Type as:  
0.##

**Check**[Show answer](#)

5) 0.625 decimal = \_\_\_\_ binary. Type as:  
0.###

**Check****Show answer**

6) 16.75 decimal = \_\_\_\_ binary. Type as:  
#####.##

**Check****Show answer**

7) 0.6 decimal = \_\_\_\_ binary, using two  
fraction bits. Type as: 0.##

**Check****Show answer**

## Binary floating-point representation

In a computer, a binary floating-point number must be represented using a fixed number of bits. Normalized scientific notation commonly using 32 or 64 bits. A common 32-bit floating-point binary representation has these items:

- Sign: 1 bit. 0 means positive, 1 negative.
- Exponent: 8 bits. Instead of two's complement, the exponent is biased, meaning a fixed value is subtracted, in this case exponent of 00000000 means  $2^0 - 127 = 2^{-127}$ , and of 11111111 means  $2^{255 - 127} = 2^{128}$ .
- Fraction: 23 bits. Because the significand's first digit is always 1, the 1 implicitly precedes the fraction and thus isn't explicitly represented. The significand in scientific notation is commonly called the *mantissa*.

### PARTICIPATION ACTIVITY

1.6.4: A 32-bit floating-point representation.

Start ☐ 2x speed

8	0	1 0 0 0 0 0 1 0	1. 0
-20	1	1 0 0 0 0 0 1 1	1. 0 1 0
31	30 29 28 27 26 25 24 23	1. 22 21 20 19 18 17 16 15	14 13 12 11 10 9 8 7 6 5 4 3 2 1 0
0 +, 1 -	$2^{(x-127)}$		
Sign	Exponent	Significand	

$$+1.0 \times 2^{(130 - 127)}$$

$$+1.0 \times 2^3 = 8$$

$$-1.25 \times 2^{(131 - 127)}$$

$$-1.25 \times 2^4 = -20$$

The above is known as the **IEEE single-precision binary floating-point** format, which uses 32 bits: 1 bit for sign, 8 bits for exponent (biased by 127), and 23 bits for significand (leading 1 before binary point is implicit). **IEEE double-precision binary floating-point** format uses 64 bits: 1 bit for sign, 11 bits for exponent, and 52 bits for significand (leading 1 before binary point is implicit).

Note: In C, C++, and Java, a variable like "float x" uses 32-bits (single precision), while "double x" uses 64 bits (double precision). Programmers usually use double to obtain more significand precision, unless memory is tightly constrained.

#### PARTICIPATION ACTIVITY

1.6.5: Representation of single-precision floating-point values.

Enter a decimal value:

Convert



### PARTICIPATION ACTIVITY

### 1.6.6: Binary floating-point to decimal.

Given the following binary floating-point representation, determine the following.

0 10000001 001000000000000000000000

1) Sign bit (type 0 or 1)

Check

Show answer

2) Exponent bits

Check

Show answer

3) Exponent in decimal

Check

Show answer



- 4) Significand's fraction bits (consider copy-pasting)

Check

Show answer

- 5) Decimal scientific notation

 × 2<sup>2</sup>

Check

Show answer

- 6) Decimal number

Check

Show answer

**PARTICIPATION  
ACTIVITY**

1.6.7: Negative exponents.

A 32-bit binary floating-point number has an exponent of 01111101 (125). Assume the sign bit is 0 and the significand is 1.000...

- 1) The number is  $+1 \times 2^?$

- ☐ 125  
☐ 127  
☐ 2  
☐ -2

- 2) The number in base ten is \_\_\_\_ .

0.25

☐ 1.000...☐ -2

3) In the 32-bit representation, the sign bit applies to \_\_\_\_ .

☐ the significand☐ the exponent☐ both the significand and the exponent

## Decimal to binary floating-point

Decimal is converted to 32-bit floating-point by first converting the decimal number to binary, then normalizing and adjusting and finally filling in the appropriate fields.

### PARTICIPATION ACTIVITY

1.6.8: Converting decimal with fraction to binary floating-point.

Start ☐ 2x speed

decimal	binary	normalized binary																								
12	1 1 0 0	1 1 0 0	x 2 <sup>3</sup>	(3 + 127 = 130)																						
2.75	1 0 . 1 1	1 0 1 1	x 2 <sup>1</sup>	(1 + 127 = 128)																						
12	0	1 0 0 0 0 0 1 0	1. 1 0																							
2.75	0	1 0 0 0 0 0 0 0	1. 0 1 1 0																							
	31	30 29 28 27 26 25 24 23	22 21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0																							
	Sign	Exponent											Significand													

**PARTICIPATION  
ACTIVITY**

## 1.6.9: Binary floating-point.

Consider converting a binary number to the above 32-bit binary floating-point format.

1) For  $1.001_2$ , the 23 fraction bits are \_\_\_\_

- ☐ 1.001000...
- ☐ 001000...

2) For  $11.101_2$ , the 23 fraction bits are \_\_\_\_ .

- ☐ 11101000...
- ☐ 1101000...
- ☐ 101000...

3) For  $1111.01_2$ , normalizing yields  $1.11101$ . Thus, the *unbiased* exponent should be \_\_\_\_ .

- ☐ 0
- ☐ 2
- ☐ 3

4) For  $10010.01$ , normalizing yields  $1.001001$ . Thus, the *biased* exponent should be \_\_\_\_ .

- ☐ 4
- ☐ 127

○ 131

**PARTICIPATION  
ACTIVITY**

## 1.6.10: Decimal to binary floating-point.

Consider converting the decimal number -16.25 to the above 32-bit binary floating-point format. Note that -16.25 is -10000.01 in binary.

10000011    4    1.000001000...    1    000001000...    131

Sign bit

Unbiased exponent in base ten

Biased exponent in base ten

Exponent bits

Fraction bits

Significand

Reset

**CHALLENGE  
ACTIVITY**

## 1.6.2: Binary floating point.

Start

Enter the decimal equivalent of the exponent of the following floating point binary number.

1 10000001 1.001000000000000000000000

Decimal exponent:

1	2	3	4	5	6
---	---	---	---	---	---

[Check](#)[Next](#)

2

3

4

5

6

Exploring further:

- [IEEE single-precision binary floating-point format](#)
- [IEEE double-precision binary floating-point format](#)

 [Provide feedback on this section](#)