

32.1 Database basics

Data

Data is numeric, textual, visual, or audio information that describes real-world systems. Data is collected and processed to aid in a variety of tasks, such as forecasting weather, analyzing financial investments, and tracking the global spread of pandemics.

Data can vary in several important ways:

- **Scope.** The amount of data produced and collected can vary. Ex: A small business might track an inventory of a few thousand items, but a large commerce website might track billions of items.
- **Format.** Data may be produced as numbers, text, image, audio, or video. Ex: A phone's proximity sensor generates raw numbers, and a satellite captures images.
- **Access.** Some data sources are private while others are made publicly available. Ex: A retail company may use private customer data to discover purchasing behavior patterns, but a government may be required by law to share certain data sets.

Historically, data was mostly **analog**, encoded as continuous variations on various physical media. Ex: Audio was recorded as vibrations impressed on vinyl disks. Images were recorded as chemicals on celluloid tapes. Today, data is mostly **digital**, encoded as zeros and ones on electronic and magnetic media.

The shift from analog to digital data facilitated the rise of large computer databases.

PARTICIPATION ACTIVITY

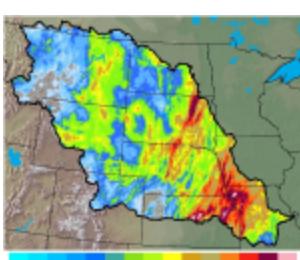
32.1.1: Examples of public data sets.



Data.gov

US Hourly Precipitation Data

Denver: 4.5, 3.4, 7.4, 3.5
Austin: 0.3, 2.1, 2.4, 0.6
Little Rock: 5.2, 6.1, 0.2, 3.5

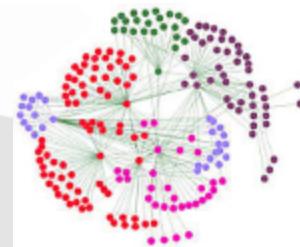


©zyBooks 01/31/24 18:17 1939727

Rob Daglio
MDCCOP2335Spring2024

Financial Tweets

@StockTwits The price of lumber was...
@TheStreet Macy's \$2M turnaround...
@MarketWatch Silbert on #bitcoin is...
@JimCramer Liked @RealMoney last....



[data.Nasa.gov](https://data.nasa.gov)

[Light Measurements](#)

9.23,9.23,9.77,9.22,8.98
10.20,7.85,8.83,9.44
6.12,8.85,9.03,8.50,1.33



©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Animation content:

Static figure:

Three public data sets appear. Each data set has a source web site, a caption describing the data set, and several rows of example data.

The first data set is from the website data.gov. The data set caption is US Hourly Precipitation Data.

The example data is:

Denver: 4.5, 3.4, 7.4, 3.5

Austin: 0.3, 2.1, 2.4, 0.6

Little Rock: 5.2, 6.1, 0.2, 3.5

The second data set is from the website kaggle.com. The caption is Financial Tweets. The example data is:

@StockTwits The price of lumber was ...

@ThStreet Macys \$2M turnaround ...

@MarketWatch Silbert on #bitcoin is ...

@JimCramer Liked @RealMoney last ...

The third data set is from the website data.nasa.gov. The caption is Light Measurements. The example data is:

9.23, 9.23, 9.77, 9.22, 8.98

10.20, 7.85, 8.83, 9.44

6.12, 8.85, 9.03, 8.50, 1.33

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Animation captions:

1. Data.gov provides thousands of U.S. government data sets. Precipitation data can be used to visualize rainfall intensity.
2. Kaggle.com allows users to find and publish data sets. The Financial Tweet data set can show who tweets on similar topics.

3. data.Nasa.gov provides data sets in aerospace and other related sciences. A data set of light measurements describes astronomical phenomena.

Sources: [Rainfall map](#), [Social network diagram](#), [Hubble mosaic of the Crab Nebula](#) from Wikipedia.org

PARTICIPATION ACTIVITY

32.1.2: Public data sets.



©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

These websites offer public data sets:

- [data.gov](#)
- [cancer.gov/research](#)
- [kaggle.com](#)
- [data.nasa.gov](#)
- [opendata.cityofnewyork.us](#)

Click on each link, review the website, and drag the website name below to the matching description.

If unable to drag and drop, refresh the page.

data.gov

cancer.gov/research

data.nasa.gov

kaggle.com

opendata.cityofnewyork.us

Provides more than 250,000 U.S. government data sets to support research and application development.

Collects and reports data and information relative to all forms of cancer.

Owned by Google, supports an online community that allows users to find and publish data sets.

Provides numerous data sets in categories such as aerospace, earth science, and space science.

Data collected by the New York City government to support continuous

monitoring and improvements to NYC
and residents' health.

Reset

Databases

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

A **database** is a collection of data in a structured format. In principle, databases can be stored on paper or even clay tablets. In practice, however, modern databases are invariably stored on computers. The database structure ensures that similar data is stored in a standardized manner.

Many modern databases contain trillions of bytes of data and support thousands of simultaneous users. Consequently, databases must be managed with sophisticated software tools:

- A **database system**, also known as a **database management system** or **DBMS**, is software that reads and writes data in a database. Database systems ensure data is secure, internally consistent, and available at all times. These functions are challenging for large databases with many users, so database systems are complex.
- A **query** is a request to retrieve or change data in a database. A **query language** is a specialized programming language, designed specifically for database systems. Query languages read and write data efficiently, and differ significantly from general-purpose languages such as Python, Java, and C++.
- A **database application** is software that helps business users interact with database systems. Many databases are complex, and most users are not familiar with query languages. Consequently, direct database access is usually not feasible. Instead, programmers write applications to simplify the user experience and ensure data access is efficient and secure.

Database software is organized in layers. Applications interact with a query language on one layer, and a query language interacts with a database system on another layer. Other software layers, such as the operating system, are beyond the scope of this material.

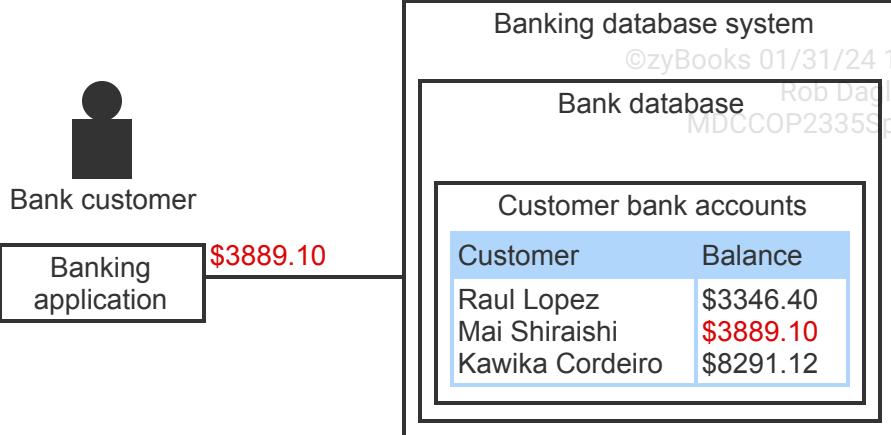
Terminology

The term **database** sometimes refers to a database system rather than the data managed by the system. The meaning is usually clear from context. In this material, **database** is used both ways.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



Animation content:

Static figure:

Three rectangles appear with captions banking database system, bank database, and customer bank account. The banking database system contains the banking database, which contains customer bank account. A table of customer names and account balances appears inside customer bank account.

A fourth rectangle has caption banking application, and is connected to the banking database system with a line.

The animation illustrates a deposit moving from the banking application to the banking database system, followed by an increase in the account balance for one customer name. The banking database system then reports the new account balance to the banking application.

Animation captions:

1. Banking data is stored in a database and is managed by a database system.
2. A bank customer uses an application to perform bank transactions.
3. Banking transactions cause the database system to modify the bank's database and update the user's account.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



Refer to the animation above.



1) Where would a bank customer's account data be directly stored?

- In the Banking database system
- In the customer banking application
- In the Bank database

2) Which of the following manages the Bank database?

- The customer banking application
- The banking database system
- The computer the Banking database system resides on

3) Which of the following would prevent unauthorized access to the Bank database?

- The banking database system
- The Bank database
- The customer banking application

4) How would a bank customer access their bank data?

- Through the banking database system
- Through the customer banking application
- Through the Bank Database



©zyBooks 01/31/24 18:17 1939727

Rob Daglio
MDCCOP2335Spring2024



Database roles

©zyBooks 01/31/24 18:17 1939727

Rob Daglio
MDCCOP2335Spring2024

People interact with databases in a variety of roles:

- A **database administrator** is responsible for securing the database system against unauthorized users. A database administrator enforces procedures for user access and database system availability.
- A **database designer** determines the format of each data element and the overall database structure. Database designers must balance several priorities, including storage, response time,

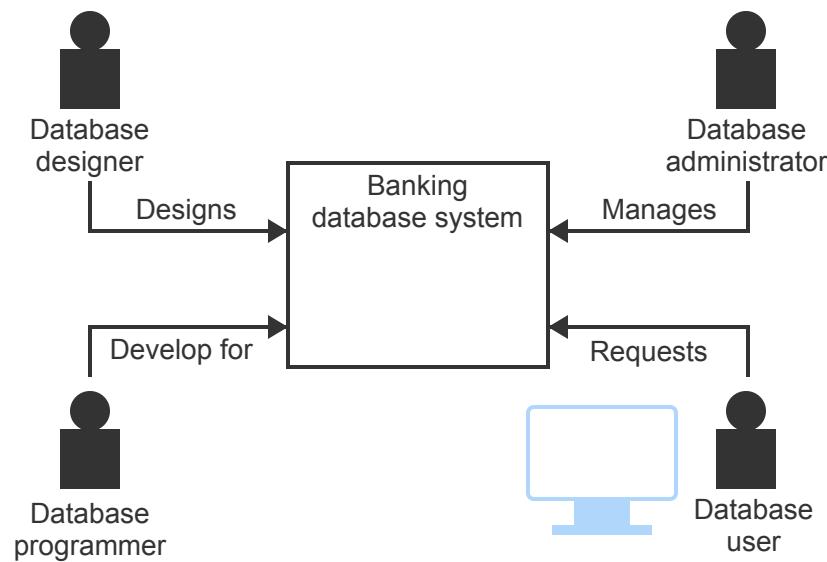
and support for rules that govern the data. Since these priorities often conflict, database design is technically challenging.

- A **database programmer** develops computer programs that utilize a database. Database programmers write applications that combine database query languages and general-purpose programming languages. Query languages and general-purpose languages have significant differences, so database programming is a specialized challenge.
- A **database user** is a consumer of data in a database. Database users request, update, or use stored data to generate reports or information. Database users usually access the database via applications but can also submit queries directly to the database system.

For simple databases with a limited amount of data and few users, one person may assume several roles. Ex: The database administrator might also be a database designer or database programmer. For large, complex databases, each person usually takes on just one role.

PARTICIPATION ACTIVITY

32.1.5: Database roles.



Animation content:

Static figure:

A box labeled banking database system appears. The box is surrounded by four icons, each representing a person and connected to the banking database system with an arrow. The four icon labels are database administrator, database designer, database programmer, and database user.

Step 1: The database designer establishes the structure of the database and determines the data to be collected and stored. The words create bank database move from the database designer to the banking database system. The words create customer and account move from the database

designer to the banking database system.

Step 2: The database administrator ensures the database is available and secure. The words add user move from the database administrator to the banking database system. The words backup data move from the database administrator to the banking database system.

Step 3: A database programmer uses query languages and programming languages to develop applications for database users. The words query database move from the database programmer to the banking database system. The words alter data move from the database programmer to the banking database system.

Step 4: Database users are the primary consumers of database data through applications and query languages. The words change address move from the database user to the banking database system. The words transfer money move from the database user to the banking database system.

Animation captions:

1. The database designer establishes the structure of the database and determines the data to be collected and stored.
2. The database administrator ensures the database is available and secure.
3. A database programmer uses query languages and programming languages to develop applications for database users.
4. Database users are the primary consumers of database data through applications and query languages.

**PARTICIPATION
ACTIVITY**

32.1.6: Roles.



1) Which role is responsible for providing access to the database?



- Database administrator
- Database designer
- Database programmer
- Database users

2) Which role is responsible for defining the detailed database design?

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



- Database administrator
- Database designer
- Database programmer
- Database users



3) Which role uses an application to query a database and generate a report?

- Database designer
- Database administrator
- Database programmer
- Database users

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Exploring further:

- [Data and Reality, by William Kent](#)
- [Data.gov](#)
- [Kaggle](#)
- [Data.NASA.gov](#)

32.2 Database systems

File systems and database systems

Small databases that are shared by one or two users can be managed in a text file or spreadsheet. Text files and spreadsheets are inadequate, however, as databases grow in size, complexity, and use. Large, complex databases that are shared by many users have special requirements:

- *Performance.* When many users and applications simultaneously access large databases, query response time degrades rapidly. Database systems maintain fast response times by structuring data properly on storage media and processing queries efficiently.
- *Authorization.* Many database users should have limited access to specific tables, columns, or rows of a database. Database systems authorize individual users to access specific data.
- *Security.* Database systems ensure authorized users only access permissible data. Database systems also protect against hackers by encrypting data and restricting access.
- *Rules.* Database systems ensure data is consistent with structural and business rules. Ex: When multiple copies of data are stored in different locations, copies must be synchronized as data is updated. Ex: When a course number appears in a student registration record, the course must exist in the course catalog.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

- Recovery. Computers, database systems, and individual transactions occasionally fail. Database systems must recover from failures and restore the database to a consistent state without loss of data.

File systems are not designed for these demanding requirements. The limitations of file systems became clear as business adopted computers in the 1960s. Since then, database systems have replaced file systems for large databases with many users.

©zyBooks 01/31/24 18:17 1939727

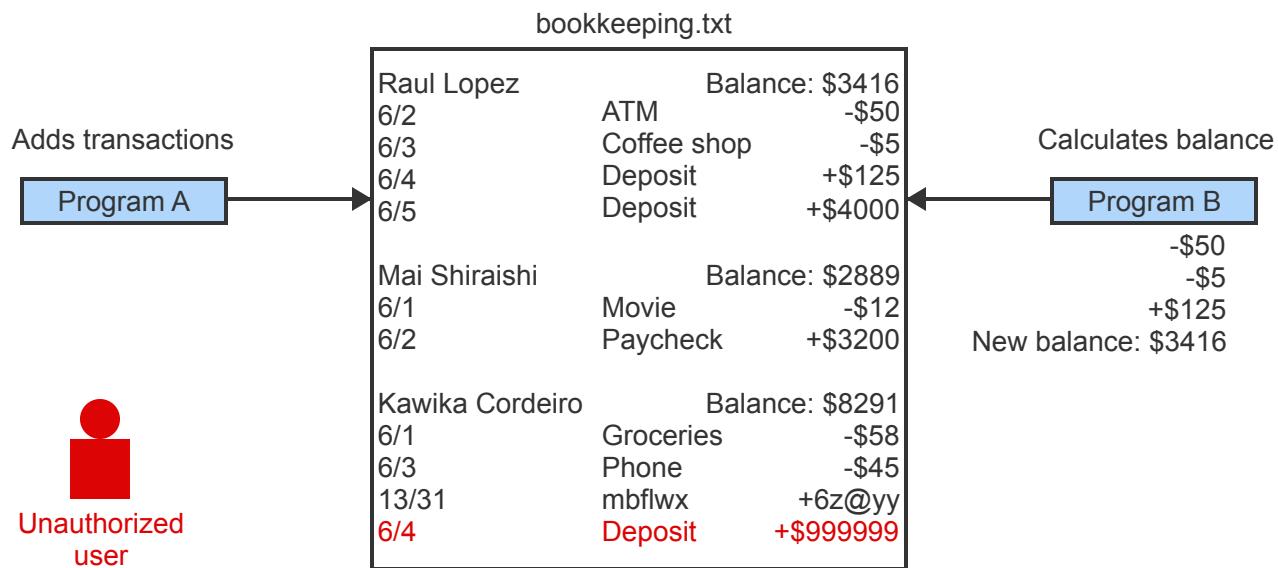
Rob Daglio

MDCCOP2335Spring2024



PARTICIPATION ACTIVITY

32.2.1: Limitations of file systems.



Animation content:

Static figure:

A file named bookkeeping.txt contains banking transaction data for three customers. For each customer, the file shows account balance and several banking transactions. Each banking transaction has a date, description, and amount.

A rectangle is named Program A with caption adds transactions. Another rectangle is named Program B with caption calculates balance. Arrows point from both rectangles to bookkeeping.txt. Off to the side, an icon of a person has caption unauthorized user.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

Step 1: A list of bank transactions is stored in the text file bookkeeping.txt. Each transaction includes a date, type, and the amount of money paid or received.

Step 2: Two programs access the text file. One adds new transactions, and the other calculates account balances.

Step 3: When program A writes transactions quickly to the file, Program B misses the \$4000 deposit and calculates Raul's balance incorrectly. Several transactions for one customer, Raul Lopez, move quickly from Program A to bookkeeping.txt. Program B calculates the new balance for Raul Lopez, but misses the last transaction and calculates the wrong balance, which is stored in bookkeeping.txt.

Step 4: Program A may write an erroneous transaction. A transaction for another customer, Kawika Cordeiro, moves from Program A to bookkeeping.txt. The transaction contains an invalid (non-numeric) amount but, nevertheless, is recorded for Kawika Cordeiro in bookkeeping.txt.

Step 5: A lack of adequate security could allow unauthorized users to access the file. A transaction moves from the unauthorized user and is recorded in bookkeeping.txt for Kawida Cordeiro.

Animation captions:

1. A list of bank transactions is stored in the text file bookkeeping.txt. Each persons' transactions include a date, type, and the amount of money paid or received.
2. Two programs access the text file. One adds new transactions, and the other calculates account balances.
3. When program A writes transactions quickly to the file, Program B misses the \$4000 deposit and calculates Raul's balance incorrectly.
4. Program A may write an erroneous transaction.
5. A lack of adequate security could allow unauthorized users to access the file.

PARTICIPATION ACTIVITY

32.2.2: Limitations of file systems.



Refer to the animation above.

- 1) If program A writes to bookkeeping.txt the same time that program B reads from bookkeeping.txt, what can potentially go wrong?



- Nothing can go wrong.
- Program A may be writing only partial data.
- Program B may be reading only partial data.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



2) Why was an unauthorized user permitted to perform the transaction that added \$999999 to bookkeeping.txt?

- The user had read access to the text file.
- The user had write access to the text file.
- The user had no access to the text file.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

3) What is wrong with the transaction posted on 13/31 for the amount +6z@yy?

- The data in the transaction contains invalid pieces of data.
- The transaction data is valid and would correctly update the bookkeeping.txt file.



Transactions

Transaction management is a particularly challenging requirement for database systems.

A **transaction** is a group of queries that must be either completed or rejected as a whole. Execution of some, but not all, queries results in inconsistent or incorrect data. Ex: A debit-credit transaction transfers funds from one bank account to another. The first query removes \$100 from one account and the second query deposits \$100 in another account. If the first query succeeds but the second fails, \$100 is mysteriously lost. The transaction must process either both queries or neither query.

When processing transactions, database systems must:

- *Ensure transactions are processed completely or not at all.* A computer or application might fail while processing a transaction. When failing to process a transaction, the database system must reverse partial results and restore the database to the values prior to the transaction.
- *Prevent conflicts between concurrent transactions.* When multiple transactions access the same data at the same time, a conflict may occur. Ex: Sam selects a seat on a flight. Maria purchases the same seat in a separate transaction before Sam completes his transaction. When Sam clicks the 'purchase' button, his seat is suddenly unavailable.
- *Ensure transaction results are never lost.* Once a transaction completes, transaction results must always be saved on storage media, regardless of application or computer failures.

©zyBooks 01/31/24 18:17 1939727

MDCCOP2335Spring2024

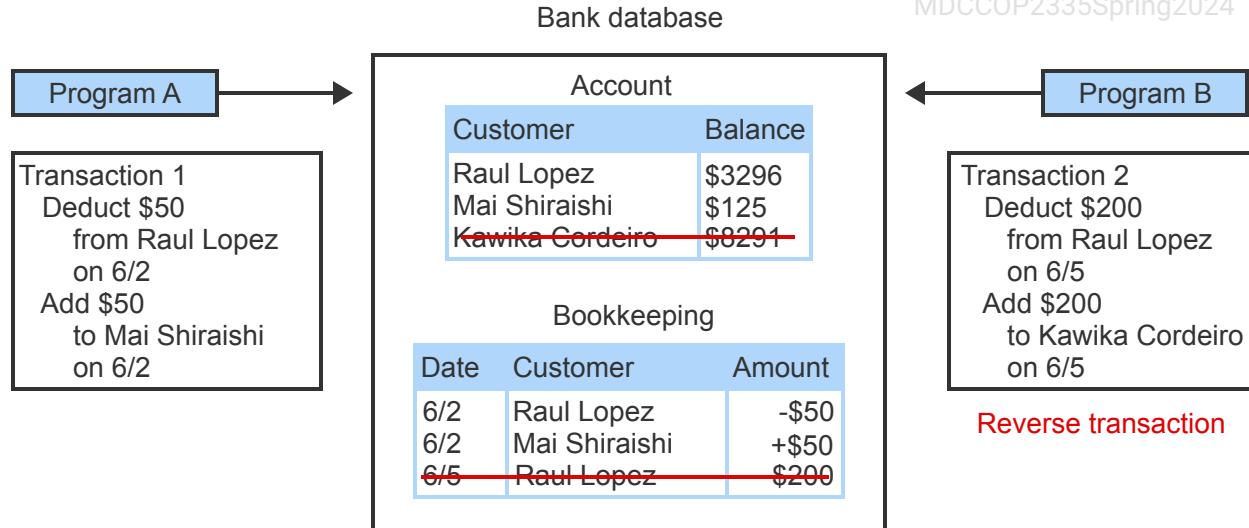
The above requirements are supported in sophisticated transaction management subsystems of most database systems.

PARTICIPATION ACTIVITY
32.2.3: Transactions.


©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



Animation content:

Static figure:

A bank database contains two tables. The first table is named Account and has columns Customer and Balance. The second table is named Bookkeeping and has columns Date, Customer, and Amount. Both tables have one row of data for Raul Lopez, one for Mai Shiraishi, and one for Kawika Cordeiro.

Below program A is a box named transaction 1 with two entries. The first entry is deduct \$50 from Raul Lopez on 6/2. The second entry is add \$50 to Mai Shiraishi on 6/2.

Below program B is a box labeled transaction 2 with two entries. The first entry is deduct \$200 from Raul Lopez on 6/5. The second entry is add \$200 to Kawika Cordeiro on 6/5. Below transaction 2 is the caption reverse transaction.

Step 1: Two programs access a bank database. The database tracks customer deposits, credits, and account balances. The account table shows entries for each of the three customers. The Bookkeeping table is empty. The transactions do not yet appear under program A and program B.

Step 2: Program A requests the database transfer \$50 from Raul to Mai. Transaction 1 appears under program A.

Step 3: Transaction 1 deducts \$50 from Raul's account and adds \$50 to Mai's account. The two entries of transaction 1 appear as rows in the Bookkeeping table. The balances in the Account table change to reflect the new rows of the Bookkeeping table.

Step 4: Program B requests that Raul transfer \$200 to Kawika. Transaction 2 appears under program B.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio
MDCCOP2335Spring2024

Step 5: \$200 is deducted from Raul's account. A third row is added to the Bookkeeping table, showing the deduction of \$200 from the Raul Lopez account. The balance for Raul Lopez is decreased by \$200 in the Account table.

Step 6: Kawika closes his account before Transaction 2 completes. The database reverses the \$200 deduction. The row for Kawika Cordeiro is crossed out in the Account table. The caption reverse transaction appears under transaction 2. The row of the Bookkeeping table, showing a deduction of \$200 from Raul Lopez account, is crossed out. The balance for Raul Lopez in the Account table is increased by \$200.

Animation captions:

1. Two programs access a bank database. The database tracks customer deposits, credits, and account balances.
2. Program A requests the database transfer \$50 from Raul to Mai.
3. Transaction 1 deducts \$50 from Raul's account and adds \$50 to Mai's account.
4. Program B requests that Raul transfer \$200 to Kawika.
5. \$200 is deducted from Raul's account.
6. Kawika closes his account before Transaction 2 completes. The database reverses the \$200 deduction.

PARTICIPATION ACTIVITY

32.2.4: Transactions.



Match the transaction behavior to the described situation.

If unable to drag and drop, refresh the page.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Transaction processed completely or not at all

Ensure transaction results are never lost

Prevent conflicts between concurrent transactions

A program is adding a penalty fee to an account that is below \$1000 while another program is adding \$2000 to the same account.

A streaming video service's database is disabled due to a lightning strike.

The IT department verifies that all database transactions performed before the lightning strike were saved on storage media.

Maria purchases an airline ticket, but a server failure causes the ticket to become unavailable before the transaction completes. The database must reverse any partial changes.

Reset

Architecture

The **architecture** of a database system describes the internal components and the relationships between components. At a high level, the components of most database systems are similar:

- The **query processor** interprets queries, creates a plan to modify the database or retrieve data, and returns query results to the application. The query processor performs **query optimization** to ensure the most efficient instructions are executed on the data.
- The **storage manager** translates the query processor instructions into low-level file-system commands that modify or retrieve data. Database sizes range from megabytes to many terabytes, so the storage manager uses **indexes** to quickly locate data.
- The **transaction manager** ensures transactions are properly executed. The transaction manager prevents conflicts between concurrent transactions. The transaction manager also restores the database to a consistent state in the event of a transaction or system failure.
- The **log** is a file containing a complete record of all inserts, updates, and deletes processed by the database. The transaction manager writes log records before applying changes to the database. In the event of a failure, the transaction manager uses log records to restore the database.
- The **catalog**, also known as a **data dictionary**, is a directory of tables, columns, indexes, and other database objects. Other components use catalog information to process and execute queries.

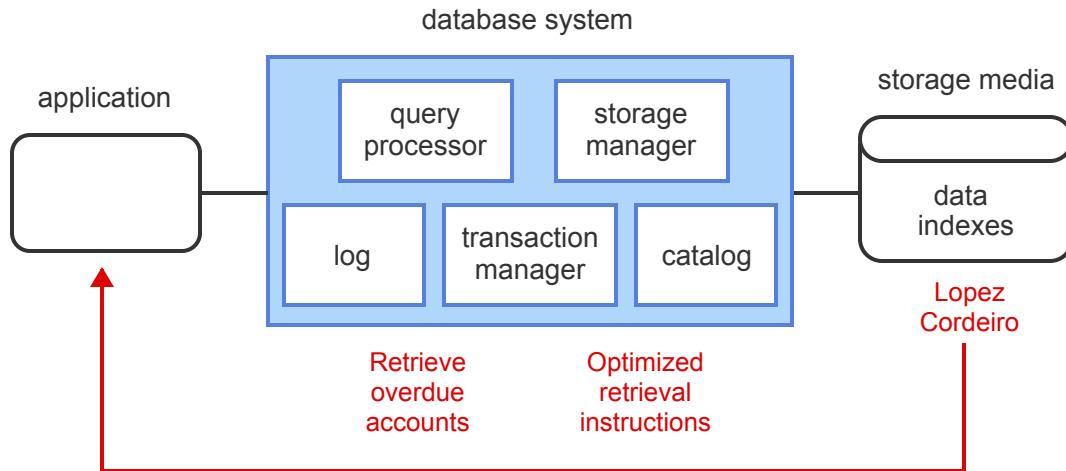
Database systems have different capabilities, and component details vary significantly. Ex: Some database systems do not support transactions and therefore have no transaction manager. Ex: The storage manager implementation depends on the physical structure of data on storage media.

PARTICIPATION ACTIVITY
32.2.5: Database system architecture.


©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



Animation content:

Static figure:

A box named database system contains boxes named transaction manager, query processor, storage manager, log, and catalog. The database system is connected by lines to a box named application and to a box named storage media.

Step 1: A database system is composed of a query processor, storage manager, transaction manager, log, and catalog. The database system and internal boxes appear.

Step 2: An application sends queries to the query processor. The application box appears. Text appears next to the application that states retrieve overdue accounts. The text moves from the application to below the database system. The query processor is highlighted.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Step 3: The query processor uses information from the catalog to perform query optimization. New text appears below the database system that states optimized retrieval instructions. The query processor and catalog are highlighted.

Step 4: The storage manager translates the query processor instructions into file system commands and uses an index to quickly locate the requested data. The storage media box appears, containing the words data indexes. The word query moves from the database system to storage media. Lopez

Cordeiro appears below storage media. The storage manager and catalog are highlighted.

Step 5: The transaction manager logs insert, update, and delete queries, and the result is sent back to the application. An arrow appears that goes from storage media to the application. The words Lopez Cordiero move along the line to the application. The transaction manager and log are highlighted.

Animation captions:

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

1. A database system is composed of a query processor, storage manager, transaction manager, log, and catalog.
2. An application sends queries to the query processor.
3. The query processor uses information from the catalog to perform query optimization.
4. The storage manager translates the query processor instructions into file system commands and uses an index to quickly locate the requested data.
5. The transaction manager logs insert, update, and delete queries, and the result is sent back to the application.

PARTICIPATION ACTIVITY

32.2.6: Database system architecture.



- 1) The query processor has direct access to the database data on storage media.



- True
 False

- 2) Without query optimization, the storage manager cannot retrieve the database data.



- True
 False

- 3) The catalog allows the storage manager to quickly locate the requested data.



- True
 False

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



4) Every database query must be logged by the transaction manager to recover the database in the event of a system failure.

- True
- False

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Products

Most leading database systems are relational. A **relational database** stores data in tables, columns, and rows, similar to a spreadsheet. All data in a column has the same format. All data in a row represents a single object, such as a person, place, product, or activity.

All relational database systems support the SQL query language. **SQL** stands for Structured Query Language and includes statements that read and write data, create and delete tables, and administer the database system.

Relational systems are ideal for databases that require an accurate record of every transaction, such as banking, airline reservation systems, and student records. The growth of the internet in the 1990s generated massive volumes of online data, called **big data**, often with poorly structured or missing information. Relational systems were not initially designed for big data and, as a result, many non-relational systems have appeared since 2000. The newer non-relational systems are called **NoSQL**, for 'not only SQL', and are optimized for big data.

Prior to 2000, most database systems were commercial products, developed by for-profit companies and licensed for a fee. Since 2000, an alternative licensing model, called open source, has become popular. **Open source** software is software that anyone can inspect, copy, and modify with no licensing fee.

NoSQL and open source systems have proliferated, and hundreds of database systems are now available. The website db-engines.com ranks systems by tracking product references on social media, internet searches, job websites, and technical websites. Internet references are an imperfect measure of product utilization, but do provide a general indication of interest and activity.

Figure 32.2.1: Leading database products.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Product	Sponsor	Type	License	rank (May 2020)
Oracle Database	Oracle	Relational	Commercial	1

MySQL	Oracle	Relational	Open source	2
SQL Server	Microsoft	Relational	Commercial	3
PostgreSQL	PostgreSQL Global Development Group	Relational	Open source	4
MongoDB	MongoDB	NoSQL	Open source	5

PARTICIPATION ACTIVITY

32.2.7: Database system categories.



Match the category to the description.

If unable to drag and drop, refresh the page.

NoSQL**Commercial****Relational****Open source**

MySQL but not MongoDB

Optimized for big data generated on the internet

SQL Server but not MySQL

Allows programmers to inspect and modify source code

Reset

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Exploring further:

- [DB-Engines Ranking](#)

32.3 Query languages

Common queries

A database system responds to queries written in a query language. A **query** is a command for a database that typically inserts new data, retrieves data, updates data, or deletes data from a database. A **query language** is a computer programming language for writing database queries.

PARTICIPATION ACTIVITY

32.3.1: Insert, select, update, and delete database queries.



Insert

Insert into Account
Ethan Carr 5000

Select

Select Name from Account
where Balance > 3000

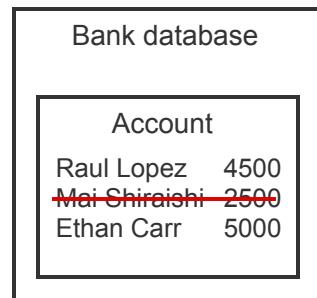
Raul Lopez
Ethan Carr

Update

Update Account
Set Raul Lopez's
balance = \$4500

Delete

Delete Mai Shiraishi
from Account



Animation content:

Step 1: A bank database stores the names and balances for two accounts: Raul and Mai. There is a box named Bank database. In this box is another box named Account. Account has the names and balances for two people.

Step 2: An insert query inserts new data into the database. Ethan's new account is inserted into the database. Text appears and states Insert into account Ethan Carr 5000. The text Ethan Carr 5000 appears in the box Account.

Step 3: A select query retrieves information from the database. The query retrieves the names of individuals that have a balance more than \$3000. New text appears and states Select name from Account where Balance is greater than 3000. The names Raul Lopez and Ethan Carr are highlighted with balances 3300 and 5000 respectively. The names appear below the text.

Step 4: An update query changes existing data in the database. Raul's balance is changed from 3300 to 4500. Text appears and states Update Account Set Raul Lopez's balance equals 4500 dollars. The number for Raul Lopez in the account box is changed from 300 to 4500.

Step 5: A delete query removes data from the database. Mai's account is removed. Text appears and states Delete Mai Shiraishi from Account. Mia Shirashi 2500 is crossed out in box Account.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Animation captions:

1. A bank database stores the names and balances for two accounts: Raul and Mai.
2. An insert query inserts new data into the database. Ethan's new account is inserted into the database.
3. A select query retrieves information from the database. The query retrieves the names of individuals that have a balance more than \$3000.
4. An update query changes existing data in the database. Raul's balance is changed from 3300 to 4500.
5. A delete query removes data from the database. Mai's account is removed.

Terminology

The four common queries are sometimes referred to as **CRUD** operations, an acronym for Create, Read, Update, and Delete data.

PARTICIPATION ACTIVITY

32.3.2: Queries.



Refer to the animation above.

- 1) Only one of the queries does not change the database contents.

- True
 False

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



2) Given the data in the Bank database, a select query for accounts with negative balances would return nothing.

- True
- False

3) The insert, select, update, and delete queries are the only types of commands necessary to interact with a database system.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

- True
- False

4) An update query cannot update data that isn't in the database.



- True
- False

Writing queries with SQL

Structured Query Language, or **SQL**, is the standard query language of relational database systems. The SQL standard is sponsored by the American National Standards Institute (ANSI) and the International Organization for Standardization (ISO). SQL is pronounced either 'S-Q-L' or 'seekwəl', but the preferred pronunciation is 'S-Q-L'.

SQL was first developed at IBM in the 1970s as an experimental query language for a prototype relational database. At the time, IBM was the dominant computer company, so SQL became the dominant relational query language. Today, all relational database systems support SQL.

Terminology

The term **NoSQL** refers to a new generation of non-relational databases. NoSQL originally meant 'does not support SQL'. However, many NoSQL databases have added support for SQL, and 'NoSQL' has come to mean 'not only SQL'.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

An SQL **statement** is a database command, such as a query that inserts, selects, updates, or deletes data:

- **INSERT** inserts rows into a table.
- **SELECT** retrieves data from a table.
- **UPDATE** modifies data in a table.
- **DELETE** deletes rows from a table.

The SQL language contains many other statements for creating and deleting databases, creating and deleting tables, assigning user permissions, and so on.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

PARTICIPATION
ACTIVITY

32.3.3: SQL statements: INSERT, SELECT, UPDATE, and DELETE.



Insert

```
INSERT INTO Account
VALUES (290, 'Ethan Carr', 5000);
```

Account		
ID	Name	Balance
831	Raul Lopez	4500
572	Mai Shiraishi	2500
290	Ethan Carr	5000

Select

```
SELECT Name
FROM Account
WHERE Balance > 3000;
```

Raul Lopez
Ethan Carr

Update

```
UPDATE Account
SET Balance = 4500
WHERE ID = 831;
```

Delete

```
DELETE FROM Account
WHERE ID = 572;
```

Animation content:

Step 1: A bank database has an Account table with three columns: ID, Name, and Balance. A new table named Account appears and has three columns named ID Name and Balance.

Step 2: The Account table has two rows that store Raul and Mai's account data. Two rows are added to the table. Row one's values are 831 Raul Lopez and 3300. Row two's values are 572 Mai Shiraishi 2500.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

Step 3: The INSERT statement adds a new row with Ethan's account data. Two lines of code appear. The first line of code states INSERT INTO Account. The second line of code states VALUES left parenthesis 290 comma apostrophe Ethan Carr apostrophe comma 5000 right parenthesis semicolon. A new row is added to the table Account with the values 290 Ethan Carr and 5000.

Step 4: The SELECT statement retrieves the names of accounts with a balance larger than \$3000. Three new lines of code appear. The first line of code states SELECT Name. the second line of code

states FROM Account. The third line of code states WHERE Balance is greater than 3000 semicolon. The values Raul Lopez and Ethan Carr are highlighted in the table Account. The values are copied and moved below the code.

Step 5: The UPDATE statement change's Raul's balance from 3300 to 4500. Three lines of code appear. The first line of code states UPDATE Account. The second line of code states SET Balance = 4500. The third line of code states WHERE ID = 831. The value in row on in column Balance int table Account changes from 3300 to 4500.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Step 6: Two lines of code appear. The first line of code states DELETE FROM Account. The second line of code states WHERE ID = 572. The second row in table Account gets a line crossed through it.

Animation captions:

1. A bank database has an Account table with three columns: ID, Name, and Balance.
2. The Account table has two rows that store Raul and Mai's account data.
3. The INSERT statement adds a new row with Ethan's account data.
4. The SELECT statement retrieves the names of accounts with a balance larger than \$3000.
5. The UPDATE statement change's Raul's balance from 3300 to 4500.
6. The DELETE statement deletes Mai from the database.

PARTICIPATION ACTIVITY

32.3.4: SQL statements.



Refer to the Account table below.

Account

ID	Name	Balance
831	Raul Lopez	3300
572	Mai Shiraishi	2500
290	Ethan Carr	5000

- 1) What is Braden Smith's balance in the following INSERT statement?



```
INSERT INTO Account
VALUES (800, 'Braden Smith',
200);
```

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

- 800
- 200
- Unknown



- 2) Which name is retrieved by the following SELECT statement?

```
SELECT Name
FROM Account
WHERE Balance < 3000;
```

- Raul Lopez
- Mai Shiraishi
- Ethan Carr

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

- 3) Whose balance does the following UPDATE statement change?

```
UPDATE Account
SET Balance = 850
WHERE ID = 290;
```

- Raul Lopez
- Mai Shiraishi
- Ethan Carr



- 4) Who is deleted by the following DELETE statement?

```
DELETE Account
WHERE ID = 999;
```

- Mai Shiraishi
- No one
- Everyone



CHALLENGE ACTIVITY

32.3.1: Query languages.



539740.3879454.qx3zqy7

Start

Country

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Name	Capital	Independence Year	Population
Syria	Damascus	1946	16910000
Namibia	Windhoek	1990	2450000

```
INSERT INTO Country
VALUES ('Lithuania', 'Vilnius', 1918, 2800000);
```

What is the new row's Population?

Pick ▼

1

2

3

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Check

Next

Creating tables with SQL

The SQL **CREATE TABLE** statement creates a new table by specifying the table and column names. Each column is assigned a **data type** that indicates the format of column values. Data types can be numeric, textual, or complex. Ex:

- INT stores integer values.
- DECIMAL stores fractional numeric values.
- VARCHAR stores textual values.
- DATE stores year, month, and day.

Some data types are followed by one or two numbers in parentheses, indicating the size of the data type. Ex: VARCHAR(10) indicates ten characters. DECIMAL(10, 3) indicates ten significant digits, including three after the decimal point.

PARTICIPATION ACTIVITY

32.3.5: Creating an Employee table.



table name
 ↓
 CREATE TABLE Employee (
 ID INT,
 Name VARCHAR(60),
 BirthDate DATE,
 Salary DECIMAL(7,2)
);

} data types

ID	Name	BirthDate	Salary
©zyBooks 01/31/24 18:17 1939727	Rob Daglio	MDCCOP2335Spring2024	

Animation content:

Static figure:

Begin SQL code:

```
CREATE TABLE Employee (
ID INT,
Name VARCHAR(60),
BirthDate DATE,
Salary DECIMAL(7,2)
);
```

End SQL code.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Step 1: The CREATE TABLE statement names the new table Employee. The first and last lines of code appear. Text appears that states table name and points to the text Employee in the first line of code. An empty table named Employee appears to the right of the code.

Step 2: The column names and data types are separated by commas. Code lines two through five appear. The words ID, Name, BirthDate, and Salary are labeled column names. The words SMALLINT, VARCHAR(60), DATE, and DECIMAL(7, 2) are labeled data types. The table Employee gets four columns named ID, Name, BirthDate, and Salary.

Animation captions:

1. The CREATE TABLE statement names the new table "Employee".
2. The column names and data types are separated by commas.

PARTICIPATION ACTIVITY

32.3.6: Creating tables with SQL.



Refer to the animation above.

- 1) The Employee table is created with 4 different data types.

- True
- False



- 2) Only the ID column stores numbers.

- True
- False

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



- 3) The BirthDate column stores only a date and no time.

- True
- False

- 4) The Employee table is an empty table once created.

- True
- False

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

PARTICIPATION ACTIVITY
32.3.7: Query the Movie table.


The SQL statements below create a Movie table and insert some movies. The SELECT statement selects all movies.

Press the Run button to produce a result table. Verify the result table displays five movies.

Modify the statement `SELECT * FROM Movie;` to only select movies released after October 31, 2015, as follows:

```
SELECT *
FROM Movie
WHERE ReleaseDate > '2015-10-31';
```

Then run the SQL again and verify the new query returns only three movies, all with release dates after October 31, 2015.

Hint: Move the semicolon that follows `FROM Movie` to the end of the statement.

```
1 CREATE TABLE Movie (
2   ID INT,
3   Title VARCHAR(100),
4   Rating VARCHAR(5),
5   ReleaseDate DATE
6 );
7
8 INSERT INTO Movie VALUES
9   (1, 'Rogue One: A Star Wars Story', 'PG-13', '2016-12-10'),
10  (2, 'Hidden Figures', 'PG', '2017-01-06'),
11  (3, 'Toy Story', 'G', '1995-11-22'),
12  (4, 'Avengers: Endgame', 'PG-13', '2019-04-26'),
13  (5, 'The Godfather', 'R', '1972-03-14');
14
15 -- Modify the SELECT statement:
16 SELECT *
17 FROM Movie;
18
```

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

[Run](#)[Reset code](#)

► View solution

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

32.4 Database design and programming

Analysis

A **database design** is a specification of database objects such as tables, columns, data types, and indexes. Database design also refers to the process used to develop the specification.

For small, simple databases, the database design process can be informal and unstructured. For large, complex databases, the process has three phases:

1. Analysis
2. Logical design
3. Physical design

The **analysis** phase specifies database requirements without regard to a specific database system. Requirements are represented as entities, relationships, and attributes. An entity is a person, place, activity, or thing. A relationship is a link between entities, and an attribute is a descriptive property of an entity.

Terminology

Analysis has many alternative names, such as conceptual design, entity-relationship modeling, and requirements definition.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Entities, relationships, and attributes are depicted in **ER diagrams**:

- Rectangles with round corners represent entities. Entity names appear at the top of rectangles.
- Lines between rectangles represent relationships.
- Text inside rectangles and below entity names represent attributes.

ER diagrams are usually supplemented by textual descriptions of entities, relationships, and attributes.

Figure 32.4.1: ER Diagram.

**PARTICIPATION ACTIVITY**

32.4.1: Analysis.



Refer to the ER diagram above.

1) What is 'Writes'?



- Entity
- Relationship
- Attribute

2) What is 'Category'?



- Entity
- Relationship
- Attribute

3) What is 'Sale'?



- Entity
- Relationship
- Attribute

4) Is there a direct relationship between Customer and Book?



- Yes
- No

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

Logical design

The **logical design** phase implements database requirements in a specific database system. For relational database systems, logical design converts entities, relationships, and attributes into tables, keys, and columns. A **key** is a column used to identify individual rows of a table. Tables, keys, and columns are specified in SQL with CREATE TABLE statements.

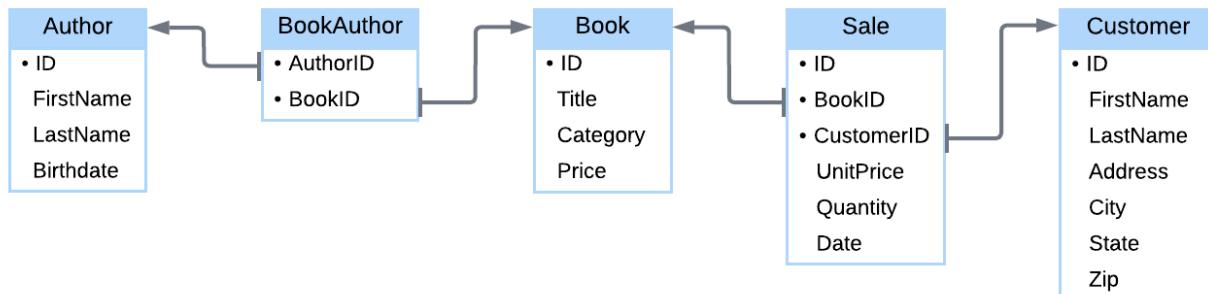
The logical design is depicted in a **table diagram**. Table diagrams are similar to ER diagrams but more detailed:

Rob Daglio
MDCCOP2335Spring2024

- Rectangles with square corners represent tables. Table names appear at the top of rectangles.
- Text within rectangles and below table names represents columns.
- Bullets (●) indicate key columns.
- Arrows between tables indicate columns that refer to keys. The tail of the arrow is aligned with the column and the arrow points to the table containing the key.

The logical design, as specified in SQL and depicted in a table diagram, is called a database **schema**.

Figure 32.4.2: Table Diagram.



PARTICIPATION ACTIVITY

32.4.2: Logical design.



Refer to the ER and table diagrams above.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

1) In the table diagram, CustomerID is:



- A table
- A key
- A column that refers to a key



2) In the table diagram, the arrow from Sale to Book corresponds to which column of Sale?

- ID
- BookID
- CustomerID

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



3) What element of the ER diagram does the BookAuthor table implement?

- Writes
- Book
- Author

4) Logical design is:

- A process only
- A specification only
- Either a process or a specification



5) A schema is depicted in:

- An ER diagram
- A table diagram



Physical design

The **physical design** phase adds indexes and specifies how tables are organized on storage media. Ex: Rows of a table may be sorted on the values of a column and stored in sort order. Physical design is specified with SQL statements such as CREATE INDEX and, like logical design, is specific to a database system.

Physical design can be depicted in diagrams. However, logical design is more important for database users and programmers, so physical design diagrams are not commonly used.

In relational databases, logical and physical design affect queries differently. Logical design affects the query result. Physical design affects query processing speed but never affects the query result. The principle that physical design never affects query results is called **data independence**.

Data independence allows database designers to tune query performance without changes to application programs. When database designers modify indexes or row order, applications run faster or slower but always generate the same results.

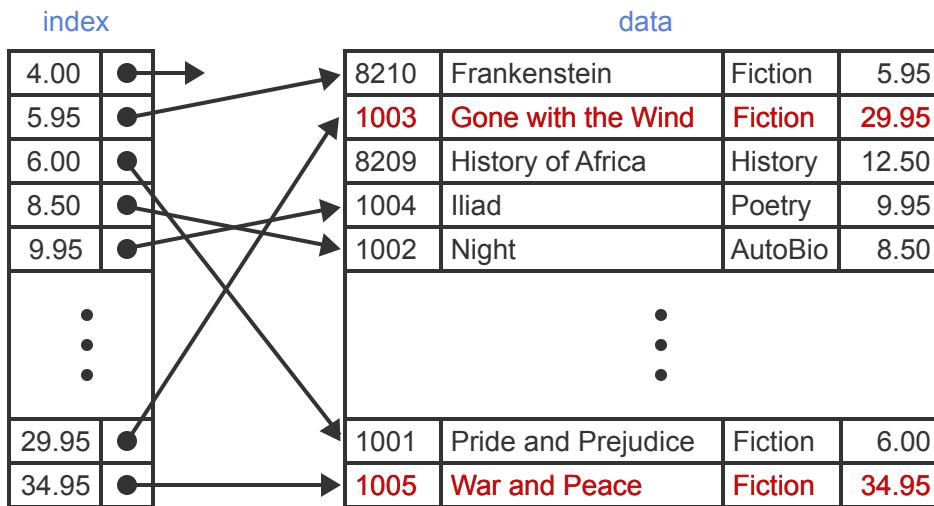
Prior to relational databases, most database systems did not support data independence. Performance tuning often forced time-consuming changes to applications. Data independence is a major advantage of relational databases and contributed to the rapid adoption of relational technology in the 1980s.

**PARTICIPATION
ACTIVITY**
32.4.3: Data independence.


©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024



```
SELECT Title
FROM Book
WHERE Price > 20.00
```

Result

Title
Gone with the Wind
War and Peace

Animation content:

Static figure:

A rectangle on the left has caption Index. A rectangle on the right has caption Data.

Index contains book prices, with two columns and seven rows. The first column has prices 4.00, 5.95, 6.00, 8.50, 9.95, 29.95, and 34.95. The second column has bullets, representing pointers to rows of Data.

Data contains book information, with four columns and seven rows. The first column contains numeric book IDs. The second column contains book titles. The third column contains book categories, such as Fiction and History. The fourth column contains book prices.

An arrow points from each bullet in Index to one row of Data. The price in the row at the arrow tail matches the price in the row at the arrowhead.

This query appears:

Begin SQL code:

```
SELECT Title
```

```
FROM Book
```

WHERE Price > 20.0;
End SQL code.

A table named Result appears below the query. The table has one column named Title. The table has two rows containing the titles War and Peace, and Gone with the Wind. In Data, the prices of these two titles are both greater than 20.00.

Step 1: The initial physical design sorts Book rows by ID. The rows of Data are sorted by the values in the ID column.

©zyBooks 01/31/24 18:17 1939727

Step 2: The SQL query selects book titles that cost more than \$20.00. Rows of Data are highlighted, one after another. The two books with price > 20.00 appear in the Result table.

Rob Daglio

MDCCOP2335Spring2024

Step 3: A new physical design sorts Book rows by Title. Rows of Data are sorted by Title.

Step 4: The new physical design has an index on Price. The Index appears, with arrows to rows of Data. Index rows are ordered by increasing price. Only the last two rows have price > 20.00.

Step 5: The query scans the index rather than the table, and thus is faster. However, Result is the same. Each row of Index is highlighted. When the last two rows are highlighted, the corresponding Data row is highlighted and the title in the row moves to Result.

Animation captions:

1. The initial physical design sorts Book rows by ID.
2. The SQL query selects book titles that cost more than \$20.00.
3. A new physical design sorts Book rows by Title.
4. The new physical design has an index on Price.
5. The query scans the index rather than the table, and thus is faster. However the result is the

PARTICIPATION ACTIVITY

32.4.4: Database design process.



Match the term to the description.

If unable to drag and drop, refresh the page.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Logical design

Database design

Physical design

Analysis

Implementation of database requirements as tables, keys, and

columns in a specific database system

The overall process of determining and implementing database requirements

Specification of database requirements without regard to implementation

@zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

Affects query performance but not query results

Reset

Programming

Because of data independence, relational database applications can be programmed before the physical design is in place. Applications may run slowly but will generate correct results.

SQL is the standard relational query language but lacks important programming features. Ex: Most SQL implementations are not object-oriented. To write a database program, SQL is usually combined with a general-purpose programming language such as C++, Java, or Python.

To simplify the use of SQL with a general-purpose language, database programs typically use an application programming interface. An **application programming interface**, or **API**, is a library of procedures or classes that links a host programming language to a database. The host language calls library procedures, which handle details such as connecting to the database, executing queries, and returning results. Ex: JDBC is a library of Java classes that access relational databases.

Dozens of database APIs are available. Each programming language supports a different API. Major programming languages like C++ and Java support several APIs.

PARTICIPATION ACTIVITY

32.4.5: Database programming with Python.



@zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

```
bookCursor = bookDatabaseConnection.cursor()
bookQuery = ('SELECT Title, Category'
             'FROM Book'
             'WHERE Price > 20.00')
bookCursor.execute(bookQuery)

for resultRow in bookCursor.fetchall():
    print('Title:', resultRow[0])
    print('Category:', resultRow[1])
```

Book				
• ID	Title	Category	Price	
1001	Pride and Prejudice	Fiction	23.00	
1002	Night	AutoBio	8.50	
1003	Gone with the Wind	Fiction	29.95	

Title: Pride and Prejudice
 Category: Fiction
 Title: Gone with the Wind
 Category: Fiction

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Animation content:

Static figure:

Begin Python code:

```
bookCursor = bookDatabaseConnection.cursor()
bookQuery = ('SELECT Title, Category'
             'FROM Book'
             'WHERE Price > 20.00')
bookCursor.execute(bookQuery)
for resultRow in bookCursor.fetchall():
    print('Title:', resultRow[0])
    print('Category:', resultRow[1])
End Python code.
```

A table named Book has columns ID, Title, Category, and Price. ID has a bullet. The table has three rows with these values:

1001, Pride and Prejudice, Fiction, 23.00

1002, Night, AutoBio, 8.50

1003, Gone with the Wind, Fiction, 29.95

A console image displays the title and category of two books:

Title: Pride and Prejudice

Category: Fiction

Title: Gone with the Wind

Category: Fiction

Step 1: The Book table contains book ID, title, category, and price.

Step 2: The Python code fragment uses the Connector/Python API to access the MySQL database system.

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Step 3: A cursor object helps extract query results. The bookCursor object connects to the book database. The first line of code is highlighted.

Step 4: bookQuery contains a SELECT query that selects the title and category for books that cost more than \$20.00. Code lines 2 through 4 are highlighted.

Step 5: The execute() method executes the SELECT query. Code line 5 is highlighted.

Step 6: Each pass through the for loop fetches one query result row into the resultRow variable. Code line 6 is highlighted.

Step 7: resultRow has an element for each result column. The print statements print both elements. The last two code lines are highlighted. The titles and categories of two books, with price > 20.00, appear on the console.

©zyBooks 01/31/24 18:17 193972
Rob Daglio
MDCCOP2335Spring2024

Animation captions:

1. The Book table contains book ID, title, category, and price.
2. The Python code fragment uses the Connector/Python API to access the MySQL database system.
3. A cursor object helps extract query results. The bookCursor object connects to the book database.
4. bookQuery contains a SELECT query that selects the title and category for books that cost more than \$20.00
5. The execute() method executes the SELECT query.
6. Each pass through the for loop fetches one query result row into the resultRow variable.
7. resultRow has an element for each result column. The print statements print both elements.

PARTICIPATION ACTIVITY

32.4.6: Database programming.



- 1) Each host language, such as Java or C++, has a different API.

- True
 False



- 2) In the animation above, the cursor helped extract individual rows of the query result for processing in a loop.

- True
 False



- 3) SQL is a general-purpose programming language.

- True
 False

©zyBooks 01/31/24 18:17 193972
Rob Daglio
MDCCOP2335Spring2024

32.5 MySQL

MySQL

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

This material uses MySQL as a reference relational database system. Although the material is relevant to all relational databases, SQL syntax and many activities are based on MySQL.

MySQL is a leading relational database system sponsored by Oracle. MySQL is relatively easy to install and use, yet has many advanced capabilities. MySQL runs on all major operating systems, including Linux, Unix, Mac OS, and Windows. For these reasons, MySQL is one of the most popular database systems.

MySQL is available in two editions:

- **MySQL Community**, commonly called **MySQL Server**, is a free edition. MySQL Server includes a complete set of database services and tools, and is suitable for non-commercial applications such as education.
- **MySQL Enterprise** is a paid edition for managing commercial databases. MySQL Enterprise includes MySQL Server and additional administrative applications.

This book is based on MySQL Server release 8.0. Forthcoming versions of this book will be upgraded to MySQL Server 8.1, released in July 2023. Complete documentation for MySQL Server 8.0 is available online.

Figure 32.5.1: MySQL documentation (dev.mysql.com/doc).

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

The world's most popular open source database 

[Contact MySQL](#) | [Login](#) | [Register](#)

MySQL™ 

MYSQL.COM DOWNLOADS DOCUMENTATION DEVELOPER ZONE

MySQL Server MySQL Enterprise Workbench InnoDB Cluster MySQL NDB Cluster Connectors More

MySQL Documentation

MySQL 8.0 Reference Manual MySQL 8.0 Release Notes Search Current Documentation 

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

Browse MySQL Documentation by: [Product](#) [Topic](#)

General	Developers & Functionality	Connectors & APIs
Getting Started with MySQL	MySQL Workbench	Connectors and APIs
Tutorial	MySQL as a Document Store	Connector/J
Server Administration	Globalization	Connector/ODBC
SQL Syntax	Optimization	Connector/.NET
InnoDB Storage Engine	Functions and Operators	Connector/Python
Alternative Storage Engines	Views and Stored Programs	PHP
Server Option / Variable Reference	Partitioning	C API
MySQL Release Notes	Precision Math	Connector/C++
MySQL Version Reference	Information Schema	MySQL for Excel
FAQs	Performance Schema	MySQL Notifier
Administrator Guides	Spatial Extensions	MySQL for Visual Studio
Installation & Upgrades	Restrictions and Limitations	NDB Cluster API Developer Guide
MySQL Server-Tool Compatibility	HA/Scalability	MySQL Enterprise
MySQL Yum Repository	MySQL InnoDB cluster	MySQL Enterprise Edition
A Quick Guide to Using the MySQL Yum Repository	MySQL NDB Cluster 8.0 (GA)	MySQL Enterprise Monitor
A Quick Guide to Using the MySQL APT Repository	MySQL NDB Cluster 7.5/7.6 (GA)	Oracle Enterprise Manager for MySQL Database
A Quick Guide to Using the MySQL SLES Repository	MySQL NDB Cluster 7.3/7.4 (GA)	MySQL Enterprise Backup
Installation Using Unbreakable Linux Network (ULN)	memcached	MySQL Enterprise Security
MySQL Installer	memcached with NDB Cluster	Secure Deployment Guide
Security	Replication	MySQL Enterprise Encryption
Secure Deployment Guide	Semisynchronous Replication	MySQL Enterprise Audit
Startup / Shutdown		MySQL Enterprise Firewall
Backup and Recovery Overview		MySQL Thread Pool
Linux/Unix Platform Guide		
Windows Platform Guide		
OS X Platform Guide		
Solaris Platform Guide		
Building from Source		
MySQL Port Reference		

Instructions for downloading and installing MySQL Server 8.0 on Windows or Mac OS are available [here](#), from the 'Exploring further' links below.

Rob Daglio
MDCCOP2335Spring2024

When installing MySQL Server, the user must enter a password for the **root account**, the administrative account that has full control of MySQL. Other database user accounts may optionally be created. After installation, MySQL Server runs as a service in the background. MySQL Server automatically starts and stops when the operating system starts and stops.

PARTICIPATION ACTIVITY

32.5.1: MySQL.



1) Refer to the website db-engines.com.



What is the overall MySQL ranking,
compared to all database systems?

- 1
- 2
- 5

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

2) What account can create other user
accounts?



- MySQL account
- Root account
- All accounts can create user
accounts

3) Content in this material only applies to
the MySQL database system.



- True
- False

MySQL Command-Line Client

The **MySQL Command-Line Client** is a text interface included in the MySQL Server download. The Command-Line Client allows developers to connect to the database server, perform administrative functions, and execute SQL statements.

To run the Command-Line Client, a user must first open a Command Prompt on Windows or a Terminal on a Mac:

- Windows: Click the Start button in the Taskbar, type "cmd", then click Command Prompt.
- Mac: Click on the Terminal application, usually found in the Applications > Utilities folder.

When MySQL Command-Line Client is started with the root account, the user is prompted to enter the root account password. Then Command-Line Client attempts to connect to the database server running on the local machine.

©zyBooks 01/31/24 18:17 1939727
MDCCOP2335Spring2024

PARTICIPATION ACTIVITY

32.5.2: Using the MySQL Command-Line Client.



```
$ mysql -u root -p
Enter password: *****
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 3
...
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
mysql> USE world;
Database changed.
mysql> SELECT * FROM city LIMIT 10;
+----+-----+-----+-----+
| ID | Name      | CountryCode | District    | Population |
+----+-----+-----+-----+
| 1  | Kabul      | AFG         | Kabul       | 1780000   |
| 2  | Qandahar   | AFG         | Qandahar   | 237500    |
| 3  | Herat      | AFG         | Herat      | 186800    |
| 4  | Mazar-e-Sharif | AFG        | Balkh      | 127800    |
| 5  | Amsterdam  | NLD         | Noord-Holland | 731200   |
| 6  | Rotterdam  | NLD         | Zuid-Holland | 593321   |
| 7  | Haag       | NLD         | Zuid-Holland | 440900   |
| 8  | Utrecht    | NLD         | Utrecht    | 234323   |
| 9  | Eindhoven  | NLD         | Noord-Brabant | 201843   |
| 10 | Tilburg    | NLD         | Noord-Brabant | 193238   |
+----+-----+-----+-----+
10 rows in set (0.00 sec)

mysql>
```

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Animation content:

Static figure:

A console appears with database commands and responses.

Step 1: From a command-line prompt, the user starts the MySQL Command-Line Client. The `-u` option names the account, and `-p` indicates a password must be entered. The following line appears on the console:

`$mysql -u root -p`

Step 2: The user enters the root password established during installation. A second line appears on the console:

Enter password: *****

©zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

Step 3: After a successful login, the user is presented a startup screen and a `mysql` prompt.

Additional lines appear on the console:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 3

...

Type help; or \h for help. Type \c to clear the current input statement.

`mysql>`

Step 4: Commands are entered one by one on the command line. The SELECT statement returns the first 10 rows from the city table. The following command appears after the mysql> prompt:
USE world;

Two more lines appear on the console:

Database changed.

mysql> SELECT * FROM city LIMIT 10;

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

A table appears on the console. The table has columns ID, Name, CountryCode, District, and Population. The table has ten rows containing data for ten cities.

Two more lines appear on the console:

10 rows in set (0.00 sec)

mysql>

Animation captions:

1. From a command-line prompt, the user starts the MySQL Command-Line Client. The -u option names the account, and -p indicates a password must be entered.
2. The user enters the root password established during installation.
3. After a successful login, the user is presented a startup screen and a mysql prompt.
4. Commands are entered one by one on the command line. The SELECT statement returns the first 10 rows from the city table.

The animation above shows the user typing SQL commands that use the 'world' database, a database that is usually installed with MySQL. The world database contains three tables: city, country, and countrylanguage. Users can practice entering SQL statements that work with and manipulate the world database. Some installations do not include the world database, so users must download and install the world database from MySQL.com separately.

MySQL Server returns an **error code** and description when an SQL statement is syntactically incorrect or the database cannot execute the statement.

Figure 32.5.2: MySQL error codes.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

```
mysql> SELECT FROM city;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to
your MySQL Server version for the right syntax to use near 'FROM city' at line 1

mysql> INSERT INTO city VALUES (123, 'Amsterdam', 'NLD', 'Noord-Holland', 731200);
ERROR 1062 (23000): Duplicate entry '123' for key 'PRIMARY'
```

@zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

PARTICIPATION ACTIVITY

32.5.3: Introduction to MySQL.



- 1) The root account password is set when installing MySQL.

- True
 False

- 2) The database server must be manually started each time the user runs the MySQL Command-Line Client.

- True
 False

- 3) The MySQL Command-Line Client provides a graphical interface for interacting with the database server.

- True
 False

MySQL Workbench

Some developers prefer to interact with MySQL Server via a graphical user interface. **MySQL Workbench** is installed with MySQL Server and allows developers to execute SQL commands using an editor. When MySQL Workbench is started, the user can connect to MySQL Server running on the local machine or on the network.

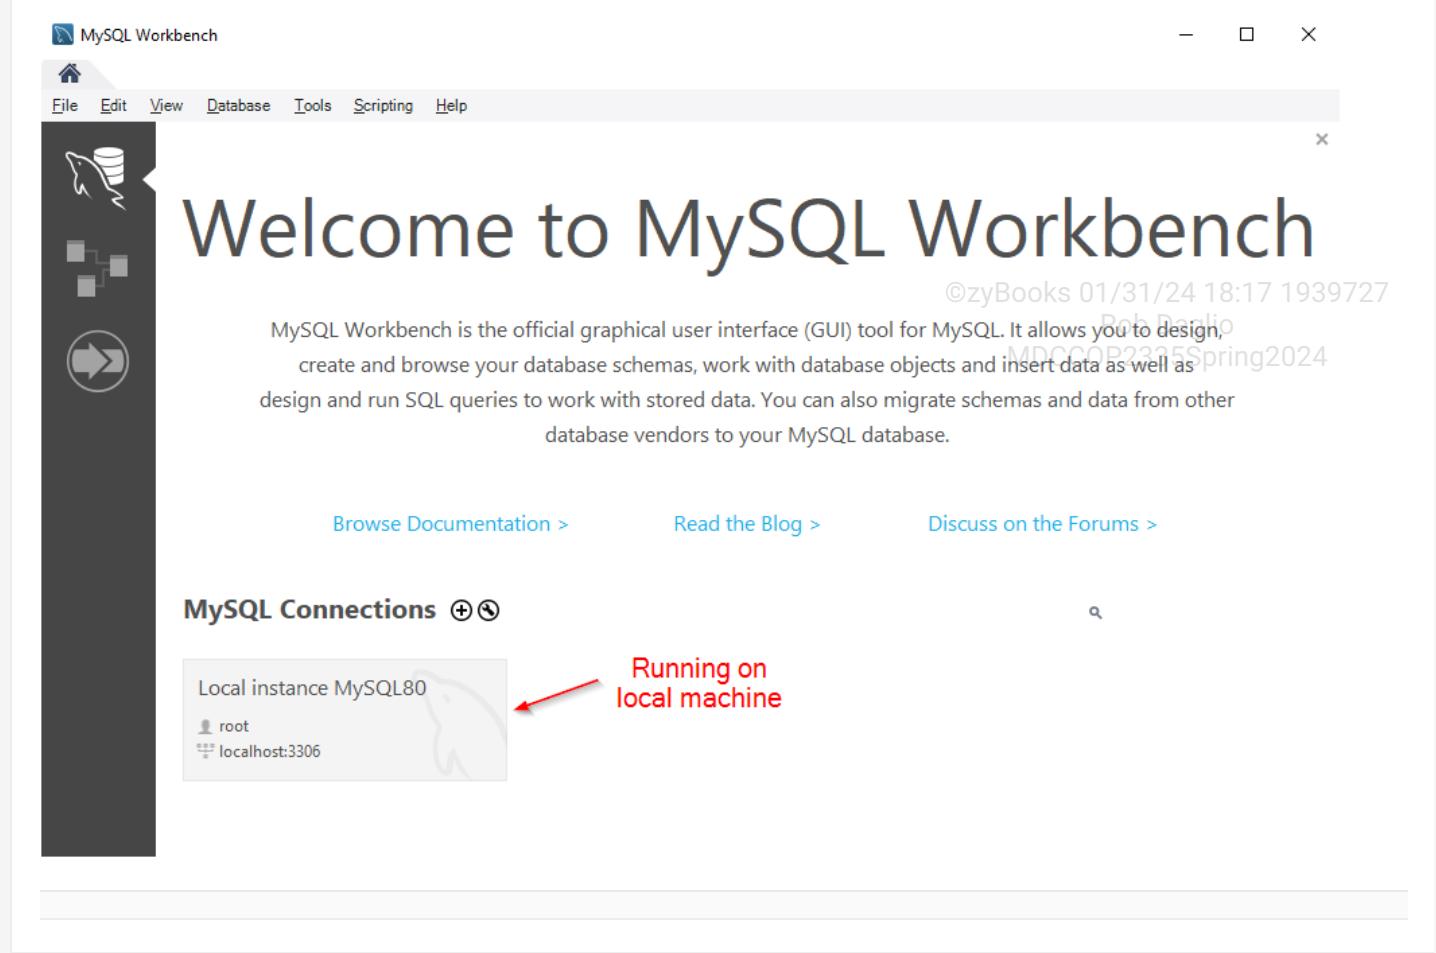
@zyBooks 01/31/24 18:17 1939727

Rob Daglio

MDCCOP2335Spring2024

The figure below shows the MySQL Workbench home screen on Windows. The Mac version has some minor differences. Clicking on the box labeled *Local Instance MySQL80* connects to MySQL Server running on the same computer as MySQL Workbench.

Figure 32.5.3: MySQL Workbench home screen.

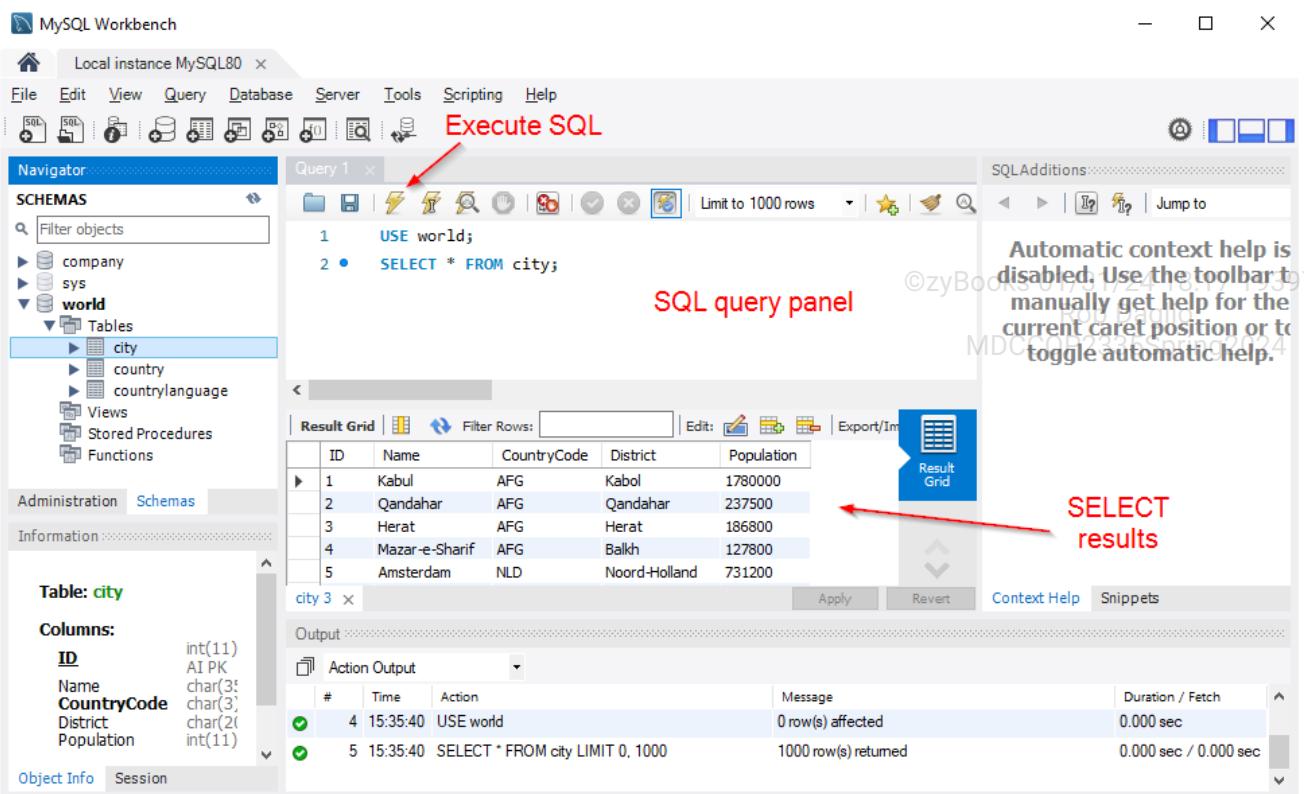


After connecting to MySQL server, Workbench shows the *Navigator* sidebar on the left with two tabs:

1. The *Administration* tab shows various administrative options, like checking the server's status, importing/exporting data, and starting/stopping the MySQL server.
2. The *Schemas* tab shows a list of available databases. A database can be expanded to show the database's tables.

The figure below shows the world database's three tables: city, country, and countrylanguage. The query panel is where the user enters SQL statements. Pressing the lightning bolt icon executes the SQL statements and shows the results below the query panel. A summary of the executed statements is shown at the bottom of the window.

Figure 32.5.4: MySQL Workbench executing a SELECT statement
©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024


PARTICIPATION ACTIVITY

32.5.4: MySQL Workbench.

- 1) MySQL Workbench and MySQL Command-Line Client both allow the user to type SQL statements.

True
 False

- 2) The SQL statements in the SQL query panel are not executed until the lightning bolt is clicked.

True
 False

- 3) The MySQL Workbench screenshot above shows the columns that make up the City table.

True
 False

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

Exploring further:

- [MySQL download](#)
- [MySQL installation video \(Windows\)](#)
- [MySQL installation video \(Mac\)](#)
- [MySQL documentation - home page](#)
- [MySQL documentation - installation](#)
- [MySQL documentation - Workbench](#)
- [MySQL documentation - problems and common errors](#)
- [MySQL documentation - Sakila database installation](#)

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

32.6 LAB - MySQL Workbench review (Sakila)

Introduction

The purpose of this lab is to gain familiarity with MySQL Workbench. The lab also ensures the correct version of the Sakila sample database is installed on your computer, for use in other zyLabs.

This lab has three parts:

- Install the Sakila database.
- Run a simple query.
- Recreate a Sakila table in the zyLab environment.

Only the third part is graded.

Install the Sakila database

This lab requires access to MySQL Server via MySQL Workbench. Most students install and access MySQL Server and MySQL Workbench on their personal computer. Installation instructions are available at [MySQL Server installation](#) and [MySQL Workbench installation](#).

To create Sakila tables in MySQL, download the [Sakila schema](#) file, open MySQL Workbench, and click the following menu commands:

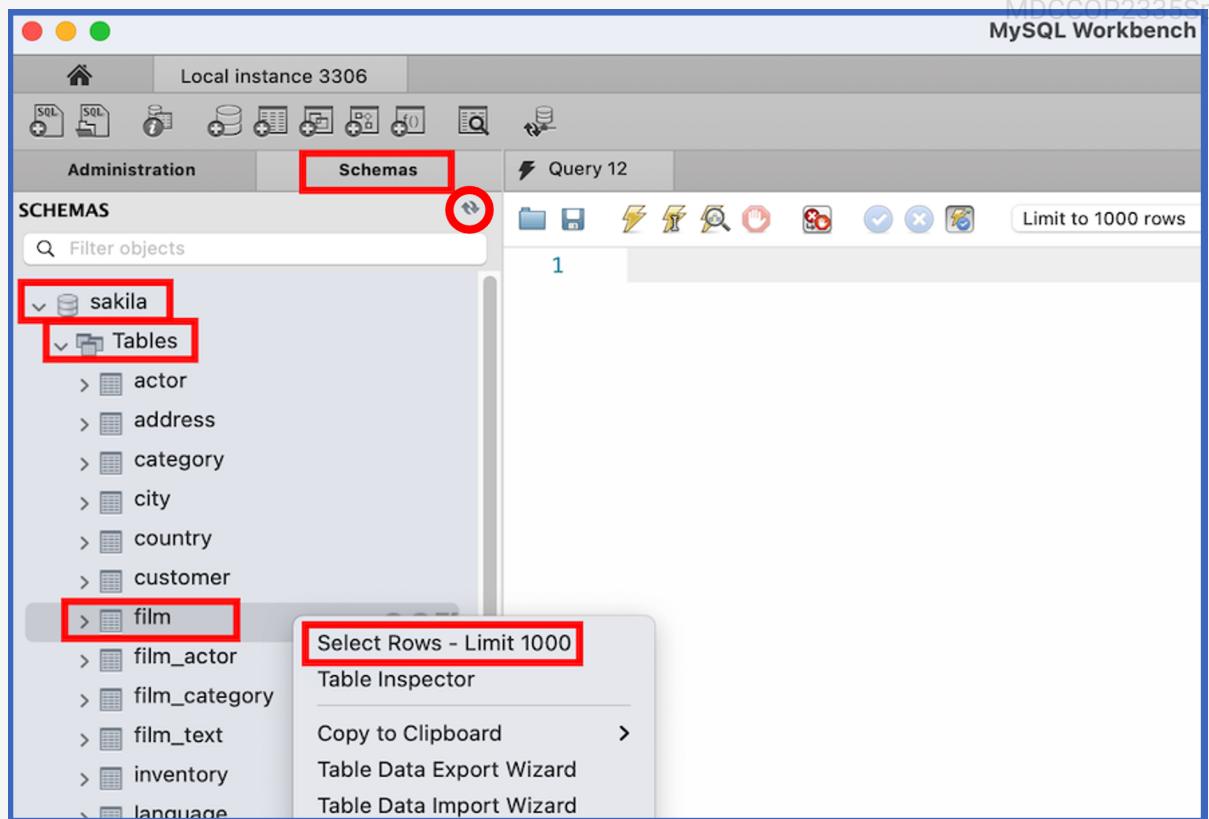
1. Click 'File' > 'Open SQL Script...' and open the Sakila schema file.
2. Click 'Query' > "Execute (All or Selection)".

To load sample data to the Sakila tables, download the [Sakila data](#) file and repeat steps 1 and 2 with this file.

Run a simple query

Refer to the following MySQL Workbench screenshot, taken from a Mac computer. Workbench looks slightly different on Windows.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024



If 'sakila' does not appear under 'Schemas', click the refresh icon, in the red circle above. If 'sakila' still does not appear, repeat the installation process or request assistance.

Depending on Workbench configuration, a different Limit may appear after 'Select Rows'.

When 'sakila' appears under 'Schemas':

1. Click > to expand 'sakila'.
2. Click > to expand 'Tables'.
3. Right-click 'film'.
4. Click 'Select Rows - Limit 1000'.

©zyBooks 01/31/24 18:17 1939727
Rob Daglio
MDCCOP2335Spring2024

MySQL Workbench executes `SELECT * FROM film;` and displays 1000 films:

The screenshot shows the MySQL Workbench interface with the Sakila database selected. In the left sidebar, the 'Tables' section is expanded, showing various tables like actor, address, category, city, country, customer, film, film_actor, film_category, film_text, inventory, language, and payment. The 'film' table is currently selected and highlighted with a gray background. In the main workspace, the SQL tab contains the query 'SELECT * FROM sakila.film;'. The results are displayed in a grid titled 'Result Grid' with columns: film_id, title, description, release_year, language_id, original_language_id, rental_duration, rental_rate, length, and replacement_cost. The first few rows of data are visible.

film_id	title	description	release_year	language_id	original_language_id	rental_duration	rental_rate	length	replacement_cost
1	ACADEMY DINOSAUR	A Epic Drama of a Feminist And a Mad Scientist...	2006	1	NULL	6	0.99	86	20.99
2	ACE GOLDFINGER	A Astounding Epistle of a Database Administrat...	2006	1	NULL	3	4.99	48	12.99
3	ADAPTATION HOLES	A Astounding Reflection of a Lumberjack And a...	2006	1	NULL	7	2.99	50	18.99
4	AFFAIR PREJUDICE	A Fanciful Documentary of a Frisbee And a Lum...	2006	1	NULL	5	2.99	117	26.99
5	AFRICAN EGG	A Fast-Paced Documentary of a Pastry Chef An...	2006	1	NULL	6	2.99	130	22.99
6	AGENT TRUMAN	A Intrepid Panorama of a Robot And a Boy who...	2006	1	NULL	3	2.99	169	17.99
7	AIRPLANE SIERRA	A Touching Saga of a Hunter And a Buller who...	2006	1	NULL	6	4.99	62	28.99
8	AIRPORT POLLOCK	A Epic Tale of a Moose And a Girl who must Co...	2006	1	NULL	6	4.99	54	15.99
9	ALABAMA DEVIL	A Thoughtful Panorama of a Database Administ...	2006	1	NULL	3	2.99	114	21.99

Recreate a Sakila table in the zyLab environment

To recreate the `actor` table in the zyLab environment:

1. Right-click 'actor'.
2. Select 'Copy to Clipboard' > 'Create Statement' to copy the `CREATE TABLE` statement to your clipboard.
3. Paste the `CREATE TABLE` statement into the zyLab Main.sql box.
4. Delete the following characters for compatibility with the zyLab environment:
 - `COLLATE=utf8mb4_0900_ai_ci`
 - all apostrophes (`)

The `CREATE TABLE` statement creates `actor` columns, keys, and indexes. No result is displayed in Develop mode. The tests in Submit mode verify that the zyLab and Sakila `actor` tables are identical.

539740.3879454.qx3zqy7

LAB ACTIVITY | 32.6.1: LAB - MySQL Workbench review (Sakila) 0 / 10

Main.sql

Load default template... ©zyBooks 01/31/24 18:17 1939727 Rob Daglio MDCCOP2335Spring2024

```

1 -- Your CREATE TABLE statement goes here
2

```

Develop mode**Submit mode**

Explore the database and run your program as often as you'd like, before submitting for grading. Click **Run program** and observe the program's output in the second box.

Run program**Main.sql**
(Your program)

Output (shown below)

Program output displayed here

Coding trail of your work

[What is this?](#)

History of your effort will appear here once you begin working on this zyLab.